

# Ordering and Consistent Cuts

Dan Deng


10/30/11

# Introduction

- Distributed systems
  - Loosely coupled processes cooperating to solve a bigger problem
- Novelty of distributed systems
  - Lamport published his paper in 1978
  - ARPANET was just “operational” in 1975
  - Temporal characteristics poorly understood
- Need a mechanism for processes to agree on time

# Distributed System Model

- Distributed system of sets of processes and channels
- Processes communicate by sending and receiving messages
- A process can observe:
  - Its own state
  - Messages it sends
  - Messages it receives
- Must enlist other processes to determine global state




# Time, Clocks, and the Ordering of events in a Distributed System – PODC influential paper award (2000)

## Leslie Lamport

(Massachusetts Computer Associates)

- B.S. in math from MIT (1960)
- Ph.D. in math from Brandeis (1972)
- Microsoft Research (2001-Current)
- Distributed systems, LaTeX
- *“a distributed system is one in which the failure of a computer you didn’t even know existed can render your own computer unusable”*



# Takeaways

- Happened-before using logical clocks to totally order events
- Logical clocks used to implement mutual exclusion
- Physical clocks for anomalous behavior

## Discussion points:

- Useful model for reasoning about temporal events
- Logical clock overflow not considered
- Does not answer precisely questions of concurrency or dependency

# Outline


- Motivation
- Partial ordering
- Logical clocks
- Total ordering
- Mutual exclusion
- Anomalous behavior
- Physical clocks

# Motivation

- Our notion of event ordering is derived using time
- Time is implemented on machines using clocks
- Local clocks on machines may not be accurate
- Need another mechanism to agree on time



# Partial Ordering



$A \rightarrow B, B \rightarrow E, A \rightarrow E$

$A \nrightarrow D, D \nrightarrow A, A \text{ and } D \text{ are concurrent}$



# Logical Clocks

- Used to implement the happened-before relation

$$A \Rightarrow B \quad C_i(A) < C_i(B)$$

- Between successive events in a process:
  - Each process increments its logical clock
- On event A of sending of a message from process  $P_i$ 
  - $P_i$  sends  $T_m = C_i(A)$  with message
- On event B of receiving of a message by process  $P_j$ 
  - B advances  $C_j(B)$  to  $\text{MAX}(T_m, C_j(B))+1$

# Total Ordering

- Happens-before gives only a partial ordering of events
- Can totally order events by
  - Ordering events by the logical times they occur
  - Break ties using an arbitrary total ordering of processes
- Specifically A happens before B if
  - $C_i(A) < C_j(B)$
  - $C_i(A) == C_j(B)$  and  $P_i < P_j$

# Total Ordering


- Can be used to solve the mutual exclusion problem in a fully distributed fashion
- Problem description:
  - Fixed number of processes
  - A single resource
  - Processes must synchronize to avoid conflict
  - Requests must be granted in order

# Mutual Exclusion

- Each process maintains its own request queue
- Process  $P_i$  – To request the resource
  - Add Request  $T_m:P_i$  to its queue
  - Send Requests  $T_m:P_i$  to all  $P_j$
- Process  $P_j$  – On receiving Request  $T_m:P_i$ 
  - Add Request  $T_m:P_i$  to its queue
  - Send Acknowledge message to  $P_i$
- Process  $P_i$  is granted resource when
  - Request  $T_m:P_i$  is earliest in request queue
  - Acknowledge is received from all  $P_j$


# Mutual Exclusion

- **Step 1:  $P_i$  Sends Request Resource**
  - $P_i$  puts **Request  $T_m:P_i$**  on its request queue
  - $P_i$  sends **Request  $T_m:P_i$**  to  $P_j$




# Mutual Exclusion

- **Step 2:  $P_j$  Adds Message**
  - $P_j$  puts **Request  $T_m:P_i$**  on its request queue
  - $P_j$  sends **Acknowledgement  $T_m:P_j$**  to  $P_i$




# Mutual Exclusion

- **Step 3:  $P_i$  Sends Release Resource**
  - $P_i$  removes **Request**  $T_m:P_i$  from request queue
  - $P_i$  sends **Release**  $T_m:P_i$  to each  $P_j$




# Mutual Exclusion

- **Step 4:  $P_j$  Removes Message**
  - $P_j$  receives **Release**  $T_m:P_i$  from  $P_i$
  - $P_j$  removes **Request**  $T_m:P_i$  from request queue





# Anomalous Behavior



- Can occur if some messages are not observed

those events in any way with the other events in  $\mathcal{L}$ , can guarantee that request A is ordered before request B.

There are two possible ways to avoid such anomalous

# Physical Clocks

- A physical clock (C) must run at about the right rate
  - $|dC_i(t) / dt - 1| < k$  where  $k \ll 1$
- Physical clocks must be somewhat synchronized
  - $|C_i(t) - C_j(t)| < \epsilon$
- Let  $\mu$  < shortest transmission time for interprocess messages
- To prevent anomalous behavior
  - $C_i(t + \mu) - C_j(t) > 0$
  - $\epsilon < \mu * (1 - k)$

# Distributed snapshots: determining global states of distributed systems

## **K. Mani Chandy** (UT-Austin)

- Indian Institute of Technology (B.E. 1965)
- Polytechnic Institute of Brooklyn (M.S. 1966)
- MIT (Ph.D. 1969)
- CS Department at UT-Austin (1970-1989)  
(department chair 1978-79 and 1983-85)
- CS Professor at CalTech (1989-Current)



## **Leslie Lamport** (SRI, 1977-1985)

# Takeaways

- Distributed algorithm to determine global state
- Detect stable conditions such as deadlock and termination
- Defines relationships among local process state, global system state, and points in a distributed computation

## Discussion points:

- Scheme accurately captures state
- Algorithm introduces communication overheads
- Related to Vector clocks



BGC Photo/Hal Korber

# Outline

- Motivation
- Distributed system model
- Consistent cuts
- Global state detection
- Stable state detection



# Motivation

- Algorithms for determining global states are incorrect
  - Relationships among local process states, global system states, and points in a distributed computation are not well understood
- Attempt to define those relationships
- Correctly identify stable states in a distributed system


# Distributed system model

- Processes
  - Defined in terms of states; states change on events
- Channels
  - State changes when messages are sent along the channel
- Events  $\mathbf{e} \langle \mathbf{p}, \mathbf{s}, \mathbf{s}', \mathbf{M}, \mathbf{c} \rangle$  defined by
  - Process  $\mathbf{P}$  in which event occurs
  - State  $\mathbf{S}$  of  $\mathbf{P}$  before event
  - State  $\mathbf{S}'$  of  $\mathbf{P}$  after event
  - Channel  $\mathbf{C}$  altered by event
  - Message  $\mathbf{M}$  sent/received along  $\mathbf{c}$



# Consistent Cuts

- Snapshot of global state in a distributed system
- Defined as snapshots where no event after the cut happened before an event before the cut
- Forbids situations where effect is seen without its cause
- Useful for debugging, deadlock detection, termination detection, and global checkpoints




# Global State Detection

- Superimposed on the computation
- Each process records its own state
- Processes of a channel cooperate on recording channel state
- Use a marker to synchronize global state recording

# Global State Detection


- Process decides to take a snapshot
  - Save its state and sends marker through its outgoing channels
  - Save messages it receives on its in channels
- Process receives a marker for the first time
  - Save state and send marker on out channels
  - Save messages it receives on its in channels
- Algorithm terminates when:
  - Each node received markers through all its incoming channels

# Global State Detection



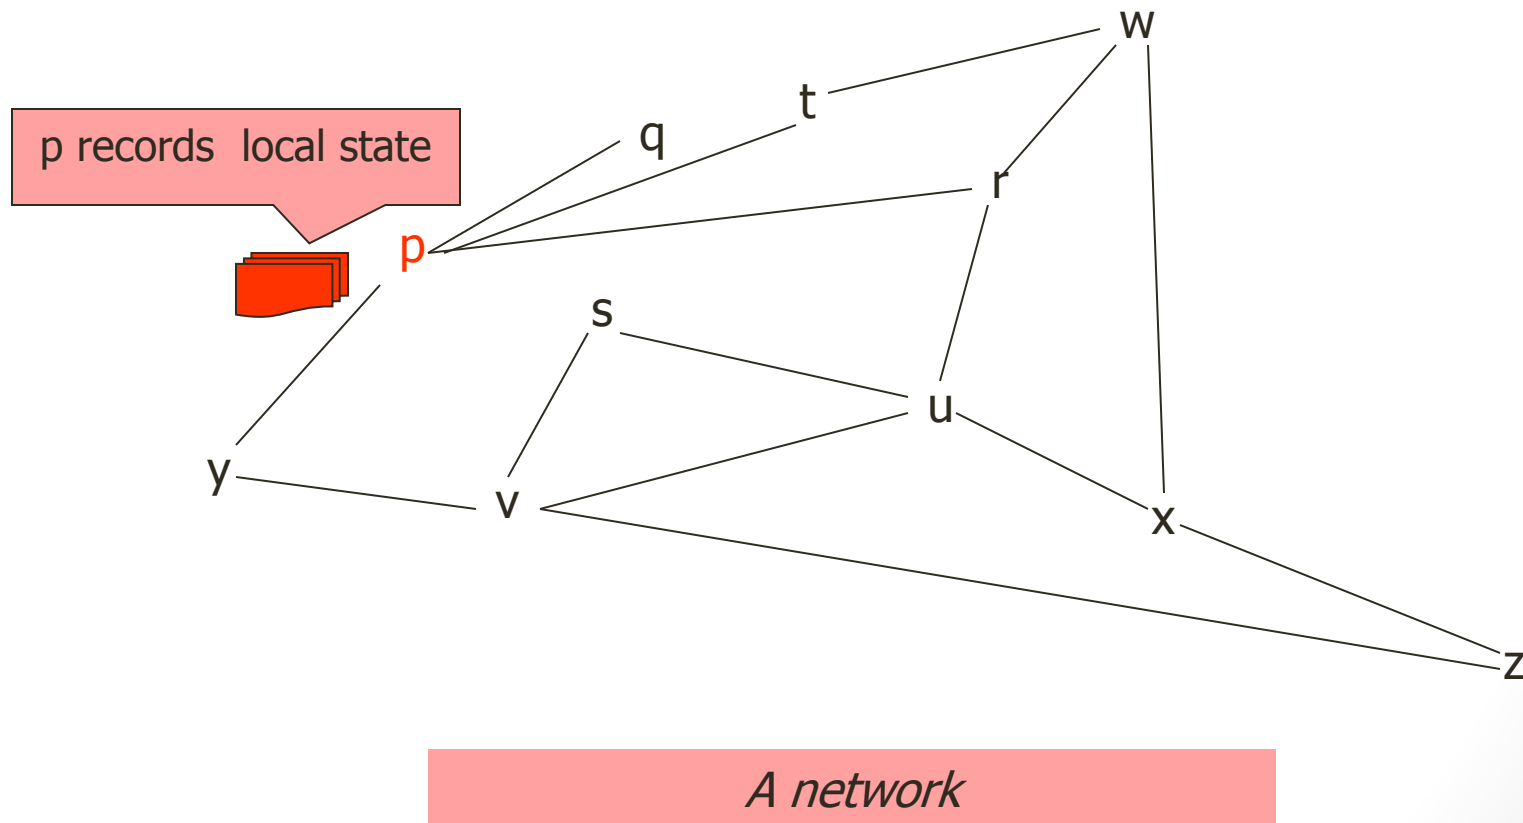
*A network*

# Global State Detection

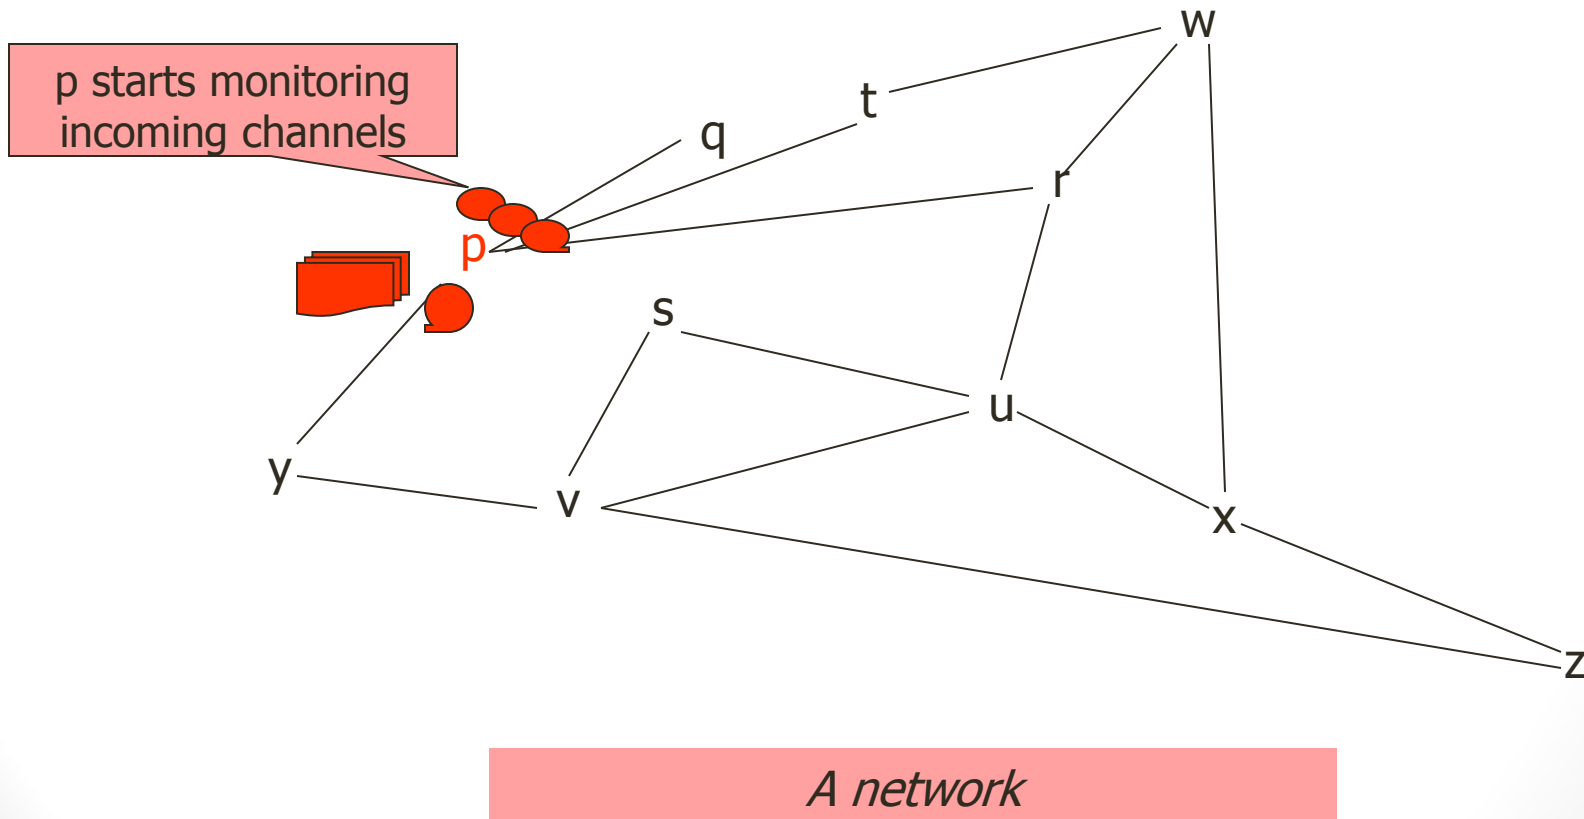


*A network*


# Global State Detection



# Global State Detection




# Global State Detection






# Global State Detection




# Global State Detection




*A network*

# Global State Detection




*A network*


# Global State Detection




# Global State Detection



# Global State Detection

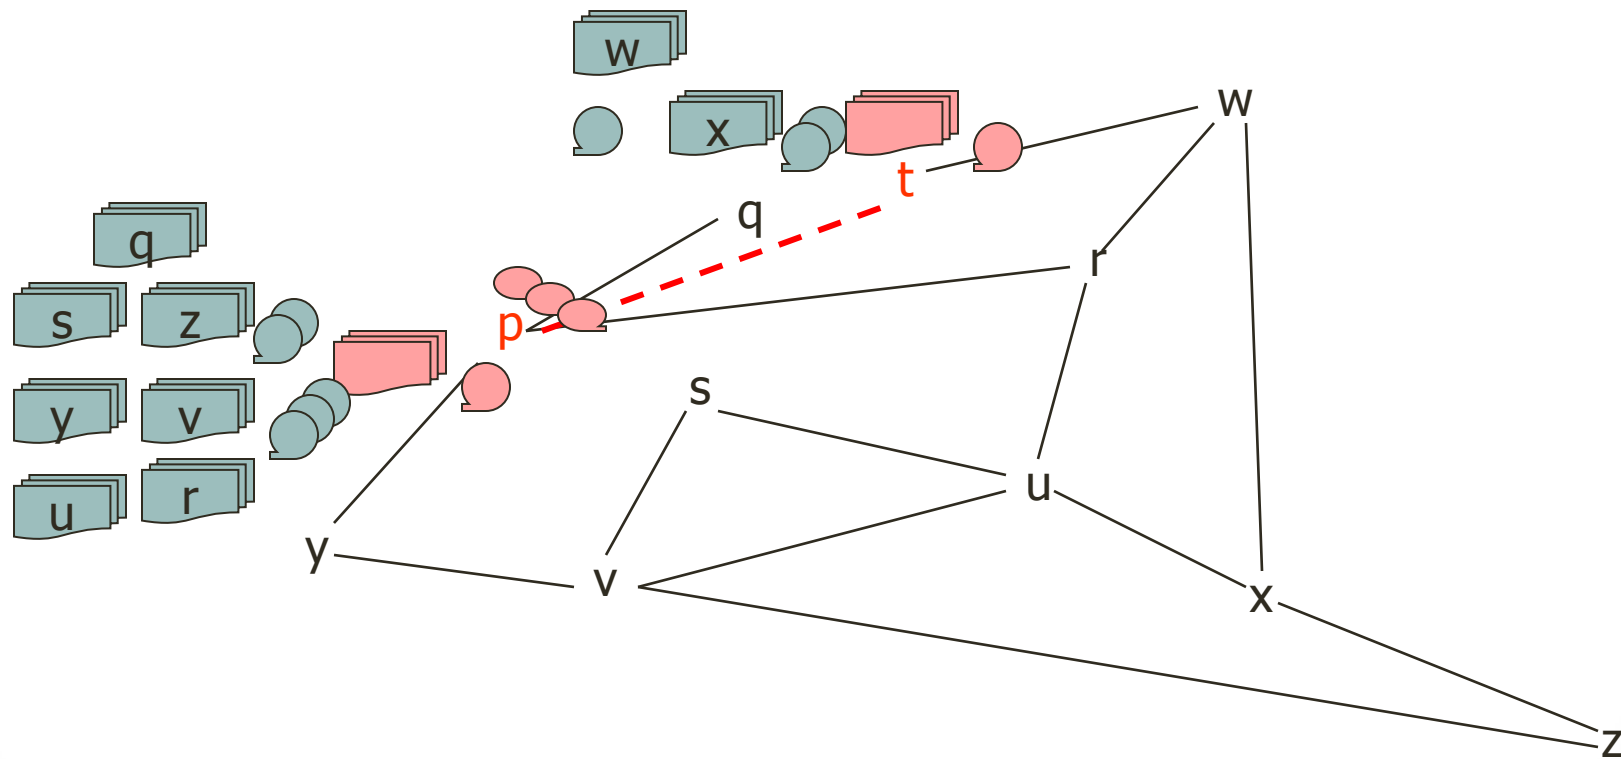


# Global State Detection



*A network*


# Global State Detection



## A network



# Global State Detection



*A snapshot of a network*

# Stable State Detection

- Input: A stable property function  $Y$
- Output: A Boolean value definite:
  - $(Y(S_i) \rightarrow \text{definite})$  and  $(Y(S_\phi) \rightarrow \text{definite})$
  - Implications of “definite”
    - $\text{definite} == \text{false}$ : cannot say YES/NO stability
    - $\text{definite} == \text{true}$ : stable property at termination
- Correctness
  - Initial state  $\rightarrow$  recorded state  $\rightarrow$  terminating state
  - for all  $j$ :  $y(S_j) = y(S_{j+1})$  – state is stable

# Takeaways

- Temporal characteristics of distributed systems was poorly understood
- Lamport proposed logical clocks for ordering
- Chandy/Lamport proposed a distributed snapshot algorithm
  - Snapshot algorithm can be used to accurately detect stable events